Evaluating AI agents

Evaluate Atlas AI agents

Copy MCP Server

Test Atlas AI agents to verify that they provide the results that you want and identify where they can improve performance.

Configure evaluations to verify that your agents provide the results that you expect. Create test cases with prompts and expected responses, run evaluations to generate agent responses, then review results to identify where your agents can improve.

Before you start

You must create an agent in the **Agent builder**.

Configure evaluation

Configure an evaluation with test cases that define prompts and expected responses.

1 **Start configuration**

1. In CDF, navigate to **Atlas AI** > **Evaluate agents**.
2. Select **+ Create new evaluation**.

2 **Enter evaluation details**

1. Enter a **Name** for your evaluation.

3 Define test cases

Add test cases to define prompts and the expected responses.

1. In the **Test cases** section, enter a **Prompt**.
2. Enter the **Expected response**. Atlas AI compares your expected response to the agent's response to test performance.
3. Optionally, choose an agent and select **Generate answer** to use the agent to generate the expected response.

> Be specific about what details, like operational status, metrics, or alert information, you want agents to include in their responses.

Run and monitor evaluation

Run an evaluation to test your agent. Monitor real-time progress as the agent processes test cases.

1 Navigate to evaluation

1. Go to **Evaluation overview**.
2. Select **Run evaluation**.

2 Select your agent

1. Use the search or filter to find your agent by name, status, or latest version.

If Atlas AI prompts you to allow code tool access for your agent, select **Confirm and run** to continue running the evaluation or **Cancel** to go back to the list of agents.

4 **Monitor evaluation progress**

When you run the evaluation, Atlas AI sends your prompts to the agent and records the responses in real time. This can take several minutes depending on how many test cases you have.

You can select **Cancel run** to stop the evaluation at any time.

> Do not close the browser tab while the evaluation is running. Evaluations run in the browser, and closing or navigating away from the tab will stop the evaluation.

View and analyze results

Review evaluation results to identify where agents need improvement and track performance across multiple runs.

1 **View completed evaluations**

On the **Evaluation overview** tab, review the list of completed evaluations.

2 **Review test case status**

Select an evaluation to view the status of your test cases.

3 **Compare responses**

builder, updating tools or data access, or editing test cases.

Learn more

- **Build and publish agents** - Learn how to create and configure agents
- **Agent tools** - Reference for available agent tools

Last modified on January 27, 2026

Was this page helpful? Yes No

About AI agent evaluations
Previous

About dashboards
Next

Cognite Docs

Cognite services